

# Coding Quiz

With the given dataset, Please compare your best possible version of

- (1) BiLSTM,
- (2) BiLSTM with multiplicative attention (you have to fix e), and
- (3) BERT

Report the accuracy, precision, recall, and f1-score of each model.

For (1) and (2), use the following hyperparameters:

Optimizer: SG  
 Embedding: GloVe (<https://pytorch.org/text/stable/vocab.html#torchtext.vocab.GloVe>) >> Please change the embed\_dim accordingly.  
 Epochs: 2  
 Batch size: 32  
 Save the model with the best params

Anything not stated, please assume accordingly

For (2), Multiplicative attention differs from the General Attention (in Assignment 4) such that, for the *Alignment Scores* (or Energy), we multiply the Keys with some weights first before we dot the Keys with the Query.

$$\mathbf{e}_i = \mathbf{q}^T \mathbf{W} \mathbf{k}_i$$

where  $\mathbf{W} \in \mathbb{R}^{h,h}$

- Hint : The shape of the Keys before and after multiplying with the weights should be the same

For (3), use this tutorial <https://huggingface.co/docs/transformers/training> (<https://huggingface.co/docs/transformers/training>) as your guide.

In [1]:

```
1 # import os
2
3 # os.environ['http_proxy'] = 'http://192.41.170.23:3128'
4 # os.environ['https_proxy'] = 'http://192.41.170.23:3128'
```

In [2]:

```

1 import torchtext
2 import torch
3 from torch import nn
4 import math
5 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
6 print(device)

```

/home/damian/pythonDSAI/lib/python3.8/site-packages/tqdm/auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See [http://ipywidgets.readthedocs.io/en/stable/user\\_install.html](http://ipywidgets.readthedocs.io/en/stable/user_install.html) (<https://ipywidget>

[s.readthedocs.io/en/stable/user\\_install.html](https://ipywidget))

from .autonotebook import tqdm as notebook\_tqdm

cuda

**1. Load the IMDB Review dataset from TorchText** (<https://pytorch.org/text/stable/datasets.html#id10>) (<https://pytorch.org/text/stable/datasets.html#id10>)

In [3]:

```

1 import torchtext
2 import torch
3 from torch import nn
4
5 #make our work comparable if restarted the kernel
6 SEED = 1234
7 torch.manual_seed(SEED)
8 torch.backends.cudnn.deterministic = True
9
10 from torchtext.datasets import IMDB
11 train_iter, test_iter = IMDB(split=('train', 'test'))
12
13
14 from torchtext.data.utils import get_tokenizer
15 tokenizer = get_tokenizer('spacy', language='en_core_web_sm')
16
17 from torchtext.vocab import build_vocab_from_iterator
18 def yield_tokens(data_iter):
19     for _, text in data_iter:
20         yield tokenizer(text)
21
22 vocab = build_vocab_from_iterator(yield_tokens(train_iter), specials=['<unk>', '<pad>'])
23 vocab.set_default_index(vocab["<unk>"])
24
25 #https://github.com/pytorch/text/issues/1350
26 from torchtext.vocab import GloVe
27 fast_vectors = GloVe(name='6B', dim=100)
28
29 fast_embedding = fast_vectors.get_vecs_by_tokens(vocab.get_itos()).to(device)
30 # vocab.get_itos() returns a list of strings (tokens), where the token at the i'th position
31 # get_vecs_by_tokens gets the pre-trained vector for each string when given a list of strings
32 # therefore pretrained_embedding is a fully "aligned" embedding matrix

```

In [4]:

```
1 input_dim = len(vocab)
2 hidden_dim = 256
3 embed_dim = 100
4 output_dim = 1
5
6 pad_idx = vocab['<pad>']
7 num_layers = 2
8 bidirectional = True
9 dropout = 0.5
10
11 batch_size = 32
12 num_epochs = 3
13 lr=0.0001
```

In [5]:

```

1 text_pipeline = lambda x: vocab(tokenizer(x))
2 label_pipeline = lambda x: 1 if x == 'pos' else 0
3
4 from torch.utils.data import DataLoader
5 from torch.nn.utils.rnn import pad_sequence #++
6
7 def collate_batch(batch):
8     label_list, text_list, length_list = [], [], []
9     for (_label, _text) in batch:
10         label_list.append(label_pipeline(_label))
11         processed_text = torch.tensor(text_pipeline(_text), dtype=torch.int64)
12         text_list.append(processed_text)
13         length_list.append(processed_text.size(0)) #++<-----packed padded sequences r
14         #criterion expects float labels
15     return torch.tensor(label_list, dtype=torch.float64), pad_sequence(text_list, padd
16
17 from torch.utils.data.dataset import random_split
18 from torchtext.data.functional import to_map_style_dataset
19
20 train_iter = IMDB(split='train')
21 test_iter = IMDB(split='test')
22
23 train_dataset = to_map_style_dataset(train_iter)
24 test_dataset = to_map_style_dataset(test_iter)
25
26 num_train = int(len(train_dataset) * 0.15)
27 num_val = int(len(train_dataset) * 0.10)
28 num_test = int(len(test_dataset) * 0.05)
29
30 split_train_, split_valid_, _ = \
31     random_split(train_dataset, [num_train, num_val, len(train_dataset) - num_train - nu
32
33 split_test_, _ = \
34     random_split(train_dataset, [num_test, len(test_dataset) - num_test])
35
36 train_loader = DataLoader(split_train_, batch_size=batch_size,
37                             shuffle=True, collate_fn=collate_batch)
38 valid_loader = DataLoader(split_valid_, batch_size=batch_size,
39                             shuffle=True, collate_fn=collate_batch)
40 test_loader = DataLoader(split_test_, batch_size=batch_size,
41                             shuffle=True, collate_fn=collate_batch)
42
43 #explicitly initialize weights for better learning
44 def initialize_weights(m):
45     if isinstance(m, nn.Linear):
46         nn.init.xavier_normal_(m.weight)
47         nn.init.zeros_(m.bias)
48     elif isinstance(m, nn.RNN):
49         for name, param in m.named_parameters():
50             if 'bias' in name:
51                 nn.init.zeros_(param)
52             elif 'weight' in name:
53                 nn.init.orthogonal_(param) #<---here
54
55 def binary_accuracy(preds, y):
56     """
57     Returns accuracy per batch, i.e. if you get 8/10 right, this returns 0.8, NOT 8
58     """
59     #round predictions to the closest integer

```

```
60 rounded_preds = torch.round(torch.sigmoid(preds))
61 correct = (rounded_preds == y).float() #convert into float for division
62 acc = correct.sum() / len(correct)
63 return acc
64
65 def train(model, loader, optimizer, criterion):
66     epoch_loss = 0
67     epoch_acc = 0
68     model.train() #useful for batchnorm and dropout
69     for i, (label, text, text_length) in enumerate(loader):
70         label = label.to(device)  #(batch_size, )
71         text = text.to(device)  #(batch_size, seq len)
72
73         #predict
74         predictions = model(text, text_length) #output by the fc is (batch_size, 1), t
75         predictions = predictions.squeeze(1)
76
77         #calculate loss
78         loss = criterion(predictions, label)
79         acc = binary_accuracy(predictions, label)
80
81         #backprop
82         optimizer.zero_grad()
83         loss.backward()
84         optimizer.step()
85
86         epoch_loss += loss.item()
87         epoch_acc += acc.item()
88
89         if i == 10:
90             break
91
92     return epoch_loss / len(loader), epoch_acc / len(loader)
93
94 def evaluate(model, loader, criterion):
95     epoch_loss = 0
96     epoch_acc = 0
97     model.eval()
98
99     with torch.no_grad():
100         for i, (label, text, text_length) in enumerate(loader):
101             label = label.to(device)  #(batch_size, )
102             text = text.to(device)  #(batch_size, seq len)
103
104             predictions = model(text, text_length)
105             predictions = predictions.squeeze(1)
106
107             loss = criterion(predictions, label)
108             acc = binary_accuracy(predictions, label)
109
110             epoch_loss += loss.item()
111             epoch_acc += acc.item()
112
113             if i == 10:
114                 break
115
116     return epoch_loss / len(loader), epoch_acc / len(loader)
117
118
```



In [6]:

```

1  class new_LSTM_cell(nn.Module):
2      def __init__(self, input_dim: int, hidden_dim: int, lstm_type: str):
3          super().__init__()
4
5          self.hidden_dim = hidden_dim
6          self.lstm_type = lstm_type
7
8          # initialise the trainable Parameters
9          self.U_i = nn.Parameter(torch.Tensor(input_dim, hidden_dim))
10         self.W_i = nn.Parameter(torch.Tensor(hidden_dim, hidden_dim))
11         self.b_i = nn.Parameter(torch.Tensor(hidden_dim))
12
13         self.U_f = nn.Parameter(torch.Tensor(input_dim, hidden_dim))
14         self.W_f = nn.Parameter(torch.Tensor(hidden_dim, hidden_dim))
15         self.b_f = nn.Parameter(torch.Tensor(hidden_dim))
16
17         self.U_g = nn.Parameter(torch.Tensor(input_dim, hidden_dim))
18         self.W_g = nn.Parameter(torch.Tensor(hidden_dim, hidden_dim))
19         self.b_g = nn.Parameter(torch.Tensor(hidden_dim))
20
21         self.U_o = nn.Parameter(torch.Tensor(input_dim, hidden_dim))
22         self.W_o = nn.Parameter(torch.Tensor(hidden_dim, hidden_dim))
23         self.b_o = nn.Parameter(torch.Tensor(hidden_dim))
24
25         if self.lstm_type == 'peephole' :
26             self.P_i = nn.Parameter(torch.Tensor(hidden_dim, hidden_dim))
27             self.P_f = nn.Parameter(torch.Tensor(hidden_dim, hidden_dim))
28             self.P_o = nn.Parameter(torch.Tensor(hidden_dim, hidden_dim))
29
30         self.init_weights()
31
32     def init_weights(self):
33         stdv = 1.0 / math.sqrt(self.hidden_dim)
34         for weight in self.parameters():
35             weight.data.uniform_(-stdv, stdv)
36
37     def forward(self, x, init_states=None):
38         bs, seq_len, _ = x.shape
39         output = []
40
41         # initialize the hidden state and cell state for the first time step
42         if init_states is None:
43             h_t = torch.zeros(bs, self.hidden_dim).to(x.device)
44             c_t = torch.zeros(bs, self.hidden_dim).to(x.device)
45         else:
46             h_t, c_t = init_states
47
48         # For each time step of the input x, do ...
49         for t in range(seq_len):
50             x_t = x[:, t, :] # get x data of time step t (SHAPE: (batch_size, input_dim, hidden_dim))
51
52             if self.lstm_type in ['vanilla', 'coupled'] :
53                 f_t = torch.sigmoid(h_t @ self.W_f + x_t @ self.U_f + self.b_f)
54                 o_t = torch.sigmoid(h_t @ self.W_o + x_t @ self.U_o + self.b_o)
55                 if self.lstm_type == 'vanilla':
56                     i_t = torch.sigmoid(h_t @ self.W_i + x_t @ self.U_i + self.b_i)
57                 if self.lstm_type == 'coupled':
58                     i_t = (1 - f_t)
59                 if self.lstm_type == 'peephole' :

```

```

60         i_t = torch.sigmoid( h_t @ self.W_i + x_t @ self.U_i + c_t @ self.P_i + self.b_i )
61         f_t = torch.sigmoid( h_t @ self.W_f + x_t @ self.U_f + c_t @ self.P_f + self.b_f )
62         o_t = torch.sigmoid( h_t @ self.W_o + x_t @ self.U_o + c_t @ self.P_o + self.b_o )
63
64         g_t = torch.tanh(          h_t @ self.W_g + x_t @ self.U_g + self.b_g )
65         c_t = (f_t * c_t) + (i_t * g_t)
66         h_t = o_t * torch.tanh(c_t)
67
68         output.append(h_t.unsqueeze(0)) # reshape h_t to (1, batch_size, hidden_dim)
69
70     output = torch.cat(output, dim = 0) # concatenate h_t of all time steps into SHAPE (seq_len, batch_size, hidden_dim)
71     output = output.transpose(0, 1).contiguous() # just transpose to SHAPE :(seq_len, batch_size, hidden_dim)
72     return output, (h_t, c_t)

```

In [7]:

```

1  import torch.nn as nn
2  from torch.nn import functional as F
3
4  class BiLSTM_model(nn.Module):
5      def __init__(self, input_dim: int, embed_dim: int, hidden_dim: int, output_dim: int):
6          super().__init__()
7          self.num_directions = 2
8          self.embedding = nn.Embedding(input_dim, embed_dim, padding_idx=pad_idx)
9          self.hidden_dim = hidden_dim
10
11         self.forward_lstm = new_LSTM_cell(embed_dim, hidden_dim, lstm_type = 'vanilla')
12         self.backward_lstm = new_LSTM_cell(embed_dim, hidden_dim, lstm_type = 'vanilla')
13
14         # These should be torch Parameters
15         self.W_h = nn.Parameter(torch.Tensor(hidden_dim*self.num_directions, hidden_dim))
16         self.b_h = nn.Parameter(torch.Tensor(hidden_dim*self.num_directions))
17
18         self.fc = nn.Linear(hidden_dim*self.num_directions, output_dim)
19
20         self.init_weights()
21
22     def init_weights(self):
23         stdv = 1.0 / math.sqrt(self.hidden_dim)
24         for weight in self.parameters():
25             weight.data.uniform_(-stdv, stdv)
26
27     def forward(self, text, text_lengths):
28         embedded = self.embedding(text)
29         embedded_flip = torch.flip(embedded, [1])
30
31         output_forward, (hn_forward, cn_forward) = self.forward_lstm(embedded, init_hidden=(0,0))
32         output_backward, (hn_backward, cn_backward) = self.backward_lstm(embedded_flip, init_hidden=(0,0))
33
34         concat_hn = torch.cat( (hn_forward, hn_backward), dim=1 )
35         ht = torch.sigmoid( concat_hn @ self.W_h + self.b_h )
36
37         return self.fc(ht)

```



In [8]:

```

1 import torch.optim as optim
2
3 bilstm = BiLSTM_model(input_dim, embed_dim, hidden_dim, output_dim).to(device)
4 bilstm.apply(initialize_weights)
5 bilstm.embedding.weight.data = fast_embedding
6
7 optimizer = optim.SGD(bilstm.parameters(), lr=lr) #<----changed to Adam
8 criterion = nn.BCEWithLogitsLoss() #combine sigmoid with binary cross entropy
9
10 best_valid_loss = float('inf')
11
12 train_losses = []
13 train_accs = []
14 # test_losses = []
15 # test_accs = []
16 valid_losses = []
17 valid_accs = []
18
19 for epoch in range(num_epochs):
20     train_loss, train_acc = train(bilstm, train_loader, optimizer, criterion)
21     # test_loss, test_acc = test(bilstm, test_loader, optimizer, criterion)
22     valid_loss, valid_acc = evaluate(bilstm, valid_loader, criterion)
23
24     train_losses.append(train_loss)
25     train_accs.append(train_acc)
26     valid_losses.append(valid_loss)
27     valid_accs.append(valid_acc)
28
29     if valid_loss < best_valid_loss:
30         best_valid_loss = valid_loss
31         torch.save(bilstm.state_dict(), 'glove_BiLSTM_attention.pt')
32
33     print(f'Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}%')
34     print(f'\t Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')
35
36 # del bilstm
37 # del optimizer
38 # del criterion

```

```

Epoch: 01 | Train Loss: 0.067 | Train Acc: 4.53%
        Val. Loss: 0.097 | Val. Acc: 7.48%
Epoch: 02 | Train Loss: 0.065 | Train Acc: 5.01%
        Val. Loss: 0.097 | Val. Acc: 7.36%
Epoch: 03 | Train Loss: 0.066 | Train Acc: 4.66%
        Val. Loss: 0.097 | Val. Acc: 7.36%


```

In [9]:

```
1 import matplotlib.pyplot as plt
2 fig = plt.figure(figsize=(10, 6))
3 ax = fig.add_subplot(1, 1, 1)
4 ax.plot(train_losses, label = 'train loss')
5 ax.plot(valid_losses, label = 'valid loss')
6 plt.legend()
7 ax.set_xlabel('updates')
8 ax.set_ylabel('loss')
```

Out[9]:

Text(0, 0.5, 'loss')




In [10]:

```
1 fig = plt.figure(figsize=(10, 6))
2 ax = fig.add_subplot(1, 1, 1)
3 ax.plot(train_accs, label = 'train acc')
4 ax.plot(valid_accs, label = 'valid acc')
5 plt.legend()
6 ax.set_xlabel('updates')
7 ax.set_ylabel('acc')
```

Out[10]:

Text(0, 0.5, 'acc')



In [11]:

```
1 bilstm.load_state_dict(torch.load('glove_BiLSTM_attention.pt'))
2 bilstm.eval()
```

Out[11]:

```
BiLSTM_model(
  (embedding): Embedding(121068, 100, padding_idx=1)
  (forward_lstm): new_LSTM_cell()
  (backward_lstm): new_LSTM_cell()
  (fc): Linear(in_features=512, out_features=1, bias=True)
)
```

In [18]:

```

1 import torch.nn as nn
2
3 #this attention mask will be apply after Q @ K^T thus the shape will be batch, seq_len,
4 def get_pad_mask(text): #[batch, seq_len]
5     batch_size, seq_len = text.size()
6     # eq(zero) is lstm output over PAD token
7     pad_mask = text.data.eq(0).unsqueeze(1) # batch_size x 1 x seq_len; we unsqueeze s
8     return pad_mask.expand(batch_size, seq_len, seq_len) # batch_size x seq_len x seq
9
10 class LSTM(nn.Module):
11     def __init__(self, len_reduction='mean'):
12         super().__init__()
13         #put padding_idx so asking the embedding layer to ignore padding
14         self.embedding = nn.Embedding(input_dim, embed_dim, padding_idx=pad_idx)
15         self.lstm = nn.LSTM(embed_dim,
16                             hidden_dim,
17                             num_layers=num_layers,
18                             bidirectional=bidirectional,
19                             dropout=dropout,
20                             batch_first=True)
21         self.fc = nn.Linear(hidden_dim * 2, output_dim)
22         self.softmax = nn.LogSoftmax(dim=1)
23         self.len_reduction = len_reduction
24         self.lin_Q = nn.Linear(hidden_dim * 2, hidden_dim * 2)
25         self.lin_K = nn.Linear(hidden_dim * 2, hidden_dim * 2)
26         self.lin_V = nn.Linear(hidden_dim * 2, hidden_dim * 2)
27
28         # lstm_output : [batch_size, seq len, n_hidden * num_directions(=2)]
29     def self_attention_net(self, lstm_output, pad_mask):
30         q = self.lin_Q(torch.clone(lstm_output))
31         k = self.lin_K(torch.clone(lstm_output))
32         v = self.lin_V(torch.clone(lstm_output))
33         # q : [batch_size, seq_len, n_hidden * num_directions(=2)]
34         # k.transpose(1, 2): [batch_size, n_hidden * num_directions(=2), seq_len]
35         # attn_w = [batch_size, seq_len, seq_len]
36
37         attn_w = torch.matmul(q, k.transpose(1, 2))
38
39         #apply padding mask
40         if self.mask:
41             attn_w.masked_fill_(pad_mask, -1e9) # Fills elements of self tensor with va
42
43         sfmt_attn_w = self.softmax(attn_w)
44         # context = [batch_size, seq_len, hidden_dim * num_directions(=2)]
45         context = torch.matmul(sfmt_attn_w, v)
46         if self.len_reduction == "mean":
47             return torch.mean(context, dim=1), sfmt_attn_w.cpu().data.numpy()
48         elif self.len_reduction == "sum":
49             return torch.sum(context, dim=1), sfmt_attn_w.cpu().data.numpy()
50         elif self.len_reduction == "last":
51             return context[:, -1, :], sfmt_attn_w.cpu().data.numpy()
52
53     def forward(self, text, text_lengths, mask=True):
54
55         self.mask = mask
56         pad_mask = get_pad_mask(text)
57
58         #text = [batch size, seq len]
59         embedded = self.embedding(text)

```

```

60
61     ### pack sequence
62     packed_embedded = nn.utils.rnn.pack_padded_sequence(embedded, text_lengths.to(
63
64     #embedded = [batch size, seq len, embed dim]
65     packed_output, (hn, cn) = self.lstm(packed_embedded) #if no h0, all zeroes
66
67     ### unpack in case we need to use it
68     output, output_lengths = nn.utils.rnn.pad_packed_sequence(packed_output, batch_
69
70     #output = [batch size, seq len, hidden dim * num directions]
71     #output over padding tokens are zero tensors
72
73     attn_output, attention = self.self_attention_net(output, pad_mask)
74
75     #attn_output = [batch_size, hidden_dim * num_direction(=2)]
76     return self.fc(attn_output), attention

```

In [19]:

```

1 #explicitly initialize weights for better learning
2 def initialize_weights(m):
3     if isinstance(m, nn.Linear):
4         nn.init.xavier_normal_(m.weight)
5         nn.init.zeros_(m.bias)
6     elif isinstance(m, nn.RNN):
7         for name, param in m.named_parameters():
8             if 'bias' in name:
9                 nn.init.zeros_(param)
10            elif 'weight' in name:
11                nn.init.orthogonal_(param) #<---here

```

In [20]:

```

1 model = LSTM().to(device)
2 model.apply(initialize_weights)
3 model.embedding.weight.data = fast_embedding ***<-----applied the fast text embedding

```

In [21]:

```

1 best_valid_loss = float('inf')
2
3 train_losses = []
4 train_accs = []
5 valid_losses = []
6 valid_accs = []
7
8 for epoch in range(num_epochs):
9     train_loss, train_acc = train(model, train_loader, optimizer, criterion)
10    valid_loss, valid_acc = evaluate(model, valid_loader, criterion)
11
12    #for plotting
13    train_losses.append(train_loss)
14    train_accs.append(train_acc)
15    valid_losses.append(valid_loss)
16    valid_accs.append(valid_acc)
17
18    if valid_loss < best_valid_loss:
19        best_valid_loss = valid_loss
20        torch.save(model.state_dict(), 'models/GloVe_BiLSTM_attention.pt')
21
22    print(f'Epoch: {epoch+1:02} | Train Loss: {train_loss:.3f} | Train Acc: {train_acc:.3f}')
23    print(f'\t Val. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')
```

AttributeError

Traceback (most recent call last)

Input In [21], in &lt;cell line: 8&gt;()

```

6 valid_accs = []
8 for epoch in range(num_epochs):
----> 9     train_loss, train_acc = train(model, train_loader, optimizer, cr
iteration)
10     valid_loss, valid_acc = evaluate(model, valid_loader, criterion)
12     #for plotting
```

Input In [5], in train(model, loader, optimizer, criterion)

```

73 #predict
74 predictions = model(text, text_length) #output by the fc is (batch_s
ize, 1), thus need to remove this 1
---> 75 predictions = predictions.squeeze(1)
77 #calculate loss
78 loss = criterion(predictions, label)
```

AttributeError: 'tuple' object has no attribute 'squeeze'

## BERT

In [22]:

```

1 import pandas as pd
2 from transformers import BertTokenizer, BertModel
3 from tqdm import tqdm
4 import numpy as np
5 from torch.optim import Adam
6
```

In [23]:

```
1 # device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
2 class BertClassifier(nn.Module):
3
4     def __init__(self, dropout=0.5):
5
6         super(BertClassifier, self).__init__()
7
8         self.bert = BertModel.from_pretrained('bert-base-cased')
9         self.dropout = nn.Dropout(dropout)
10        self.linear = nn.Linear(768, 5)
11        self.relu = nn.ReLU()
12
13    def forward(self, input_id, mask):
14
15        _, pooled_output = self.bert(input_ids= input_id, attention_mask=mask, return_dict=True)
16        dropout_output = self.dropout(pooled_output)
17        linear_output = self.linear(dropout_output)
18        final_layer = self.relu(linear_output)
19
20    return final_layer
```



In [27]:

```

1  def trainb(model, train_loader, valid_loader, learning_rate, epochs):
2
3      device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
4
5      criterion = nn.CrossEntropyLoss()
6      optimizer = Adam(model.parameters(), lr= learning_rate)
7
8      if device:
9          model = model.cuda()
10         criterion = criterion.cuda()
11
12     for epoch_num in range(epochs):
13
14         total_acc_train = 0
15         total_loss_train = 0
16
17         for train_input, train_label in tqdm(train_loader):
18
19             train_label = train_label.to(device)
20             mask = train_input['attention_mask'].to(device)
21             input_id = train_input['input_ids'].squeeze(1).to(device)
22
23             output = model(input_id, mask)
24
25             batch_loss = criterion(output, train_label)
26             total_loss_train += batch_loss.item()
27
28             acc = (output.argmax(dim=1) == train_label).sum().item()
29             total_acc_train += acc
30
31             model.zero_grad()
32             batch_loss.backward()
33             optimizer.step()
34
35         total_acc_val = 0
36         total_loss_val = 0
37
38         with torch.no_grad():
39
40             for val_input, val_label in val_loader:
41
42                 val_label = val_label.to(device)
43                 mask = val_input['attention_mask'].to(device)
44                 input_id = val_input['input_ids'].squeeze(1).to(device)
45
46                 output = model(input_id, mask)
47
48                 batch_loss = criterion(output, val_label)
49                 total_loss_val += batch_loss.item()
50
51                 acc = (output.argmax(dim=1) == val_label).sum().item()
52                 total_acc_val += acc
53
54         print(
55             f'Epochs: {epoch_num + 1} | Train Loss: {total_loss_train / len(train_loader)} | '
56             f'Val Loss: {total_loss_val / len(valid_loader)} | '
57             f'Acc: {total_acc_val / len(valid_loader)}')
58
59     return total_acc_val

```

In [28]:

```

1 EPOCHS = 5
2 model = BertClassifier()
3 LR = 1e-6
4
5 trainb(model, train_loader, valid_loader, LR, EPOCHS)

```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertModel: ['cls.seq\_relationship.weight', 'cls.seq\_relationship.bias', 'cls.predictions.bias', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.LayerNorm.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).  
 - This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

0% | 0/118 [00:00<?, ?it/s]

-----  
 ValueError

Traceback (most recent call last)

Input In [28], in <cell line: 5>()

2 model = BertClassifier()

3 LR = 1e-6

----> 5 trainb(model, train\_loader, valid\_loader, LR, EPOCHS)

Input In [27], in trainb(model, train\_loader, valid\_loader, learning\_rate, epochs)

14 total\_acc\_train = 0

15 total\_loss\_train = 0

----> 17 for train\_input, train\_label in tqdm(train\_loader):

19 train\_label = train\_label.to(device)

20 mask = train\_input['attention\_mask'].to(device)

ValueError: too many values to unpack (expected 2)

In [ ]:

1

In [ ]:

1